

Concept-C#

(Type)Classes for the Masses

Claudio Russo Matt Windsor

Friday 13 November 2017

Concepts are a proposed way to structure polymorphic C# code. Like interfaces, they constrain generic type arguments based on the presence of certain functions over those types. Unlike interfaces, but like extension methods, concepts separate the implementation of the functions from the types themselves, and do not tie the functions to existing objects. This lets programmers connect existing types to concepts, and abstract over static methods like constructors and type-level properties.

Introduction

NB: This document is still being written.

When writing generic code, we often rely on a handful of facts about the types used with that code. Most modern languages include *bounded parametric polymorphism*, or the ability to do this ahead of time in a type-safe manner¹. There are two main approaches:

- *subtype polymorphism*: a bound on type T tells us that T is a subtype of some other type U , giving generic code access to all of the methods of, or able to accept a, U^2 ;
- *typeclasses*: a bound on type T tells us that T is a member of some group (a ‘class’) C , giving generic code access to any methods that are part of the definition of, or able to accept members of, C .

Both have the same idea—they assign functionality to types—but different strengths and weaknesses. Subtype polymorphism is a natural fit for object-oriented languages, but bakes whether or not a type satisfies certain bounds indelibly into that type’s definition. Typeclasses are more flexible, and scale easily to bounds relating multiple types in complex ways, but are more difficult to weave into an object-oriented system.

.NET, Microsoft’s ecosystem for managed code, is object-oriented at its core. Since .NET 2.0, it has had a high-quality generic type system with subtype polymorphism based

¹Except Go.

²If the polymorphism system supports variance, sometimes the bound is flipped: T is a U -supertype of U .

on classes—bundles of data and methods to operate on them—and *interfaces*—packs of method signatures that can be treated as subtypes for polymorphism. Interfaces deal with many of the same problems in bounded polymorphism as typeclasses, but inherit³ many limitations from subtype polymorphism:

- They can only constrain the members of one type—their subtype.
- In .NET, they cannot abstract over constructors or other static members.
- It is hard to use the actual eventual type of an interface implementation in the interface method specification. This usually results in unnecessary *boxing* of said type into its interface form, which affects performance.

In fact, modern object-oriented programming languages are increasingly experimenting with a move towards typeclass-based bounded polymorphism. Languages like *Rust* have adopted typeclasses as the core of their generic type system. Other languages, like *Swift*, have hybrid systems, taking cues from both interfaces and typeclasses. At least one language, *Scala*, has both interfaces *and* typeclasses!

With this in mind, we propose *concepts*, a C# implementation of typeclass-style bounded parametric polymorphism. At first, this model seems alien to C#. However, as a recent language design proposal suggested, concepts find a home in the existing C# model as an extension of its existing *extension methods* system. Not only do concepts help us gain the expressiveness of powerful typeclass based bounded polymorphism, but they also let us tap into the specialisation power of modern .NET just-in-time compilers to give performance rivalling hand-specialised code. To assess the impact of concepts, we have a fork of the *Roslyn* C# compiler, which we call *Concept-C#*.

This document

In this document, we:

- Look at the key parts of the C# type system, including extension methods;
- Give our prototype design for concepts by example, starting with the basics and proceeding to more exotic features available in the *Concept-C#* prototype;
- Discuss how we implement *Concept-C#* on top of the existing Common Language Runtime and its subtype-based generics system;
- Conclude by discussing related and future work.

The C# 7.1 type system

As one of the key languages of the .NET ecosystem, C#'s type system is heavily based on the underlying .NET Common Language Runtime (CLR). The CLR supports generic methods and types, with a rich system of constraints, variance, and subtyping.

³Pun intended.

Classes

Following in the Simula tradition, classes contain all data and code inside a C# program. While most classes must be instantiated as objects to be used, C# also supports static members, which can be used without an object reference.

Following in another tradition, we begin exploring C#'s object orientation by looking at how we might store shapes in C#. A Rectangle class might look as follows:

```
public class Rectangle
{
    // Data (fields)
    private int _length;
    private int _breadth;

    // Constructor
    Rectangle(int length, int breadth)
    {
        _length = length;
        _breadth = breadth;
    }

    // Instance method
    public int Area()
    {
        return _length * _breadth;
    }

    // Properties
    public int Length => _length;
    public int Breadth => _breadth;

    // Static method
    public static Rectangle Square(int length)
    {
        return new Rectangle(length, length);
    }
}
```

Interfaces

If we have a method that works with the area of Rectangles, we may also need it to handle the area of Triangles, Circles, and other shapes. If each defines a method

```
public int Area();
```

then we can call upon C#'s interface-based subtyping system as follows:

```
public interface Shape
{
    int Area();
}
```

```
public class Rectangle : Shape
{
    // ...
    public int Area()
    {
        return _length * _breadth;
    }
}
```

```
public class RightAngledTriangle : Shape
{
    // ...
    public int Area()
    {
        return (_opposite * _adjacent) / 2;
    }
}
```

The subtyping relationship between these classes is manifest in the definitions of the types themselves. This means we run into problems if, say, we use a library that contains a nice implementation of `Circle`, but didn't have the foresight to implement our `Shape` interface.

Generics

In the C# syntax, we mark generic methods and types with an angle-bracket-delimited *type parameter list*. To use generics, we *construct* the generic entity with a mapping of type arguments to type parameters. For example, the .NET Framework ships with a generic class, `Lazy<T>`, which lets a value of type `T` be lazily initialised when, and only when, its actual value is needed. `Lazy<int>` is a possible construction of `Lazy<T>` for lazy initialisation of integer values.

Each type parameter corresponds to an extra obligation when building a type deduction for its generic. For example, we deduce `Lazy<int>` as follows:

$$\frac{T \vdash \text{int}}{T \vdash \text{Lazy} \langle \text{int} \rangle} \text{Lazy}$$

Operator overloading and generic arithmetic

When we calculate the area of our Shapes, we use C#'s builtin operators `*` and `/` to multiply and divide integers. Perhaps, however, we need to express the dimensions of our Shapes in some other type. The builtin operators have us covered if we move to floating point numbers, or some other size of integer—but what of `Lazy<int>`, or any custom integer type we make or stumble across in the future?

To let programmers define their own numeric types, and use the familiar numeric notation, C# allows for *operator overloading*. Classes can define static methods with certain syntax, and C# will call into them when it encounters the corresponding operator.

This is useful as a way of making numeric code cleaner and easier to read, but can we abstract over the presence of such operators? Can we make a `Shape<TArea>`, where `TArea` supports `*` and `/`?

Alas, no: operator overloads sit outside the subtyping system! To abstract over generic arithmetic, we would need to create a new interface, at which point:

- We can't make our arithmetic generic over existing classes, as they can't make themselves subtypes of the interface;
- We can't make our arithmetic generic over primitive types, as they can't subtype at all;
- Even if we restrict ourselves to our own classes, we can't represent all of the operations we would need for practical generic arithmetic as interface methods.

Let's explore the third point. Most types that act like numbers work with a common set of operations: addition, subtraction, multiplication, splitting into absolute value and sign, and conversion from an integer⁴. We could try to make this set into a C# interface, but quickly run into trouble:

```
public interface INum<T>
{
    T Add(T other);
    T Sub(T other);
    T Mul(T other);
    T Abs();
    T Signum();
    // ...we can't implement conversion from integer:
    // we don't have an INum<T> to call it from!
}
```

⁴We discuss this set in particular as it corresponds to Haskell's `Num` typeclass.

```
// How do we make 'int' an INum<int>?
```

Extension methods

A key problem of subtype-based systems is that they are closed for extension. If we want to add a class to an interface, we have to modify the class's code. If we don't own the code we need to extend, we must resort to workarounds like static methods and type wrappers.

This problem extends to adding new methods to existing types. For example, maybe we really need to add the ability to rotate a `Rectangle` 90 degrees, but aren't allowed to add a new `Rotate` method to the class itself. To fix this, C# 3.0 introduced *extension methods*—static methods that can be invoked as if they were native methods on some target type:

```
public static class ShapeExtensions
{
    public static Rectangle Rotate(this Rectangle me)
    {
        return new Rectangle(length: me.Breadth, breadth: me.Length)
    }
}
```

```
var rec = new Rectangle(27, 53); // length 27, breadth 53
var rot = rec.Rotate();           // breadth 27, length 53
```

Extension methods are a powerful way to extend existing classes, but do have limitations. We can't constrain type parameters based on whether they have certain extension methods available, and certainly can't implement interfaces using them.

Concepts

In our tour of C# 7.1, we found some pain points with subtyping polymorphism:

- We can't add interfaces to existing types, including primitive types;
- We can't fit static methods, type-level properties, operator overloads, and other such exotica inside interfaces;
- Extension methods let us add functionality to existing types, but we can't use them in bounds;

In this section, we explore our design for concepts in C#, showing how concepts can address all of these points.

Basic concepts

A concept names several methods and properties, and groups together types based on whether an implementation exists for each. Unlike interfaces, the methods and properties exist *at the type level*—they are not called on objects.

To underline the difference, here is a concept version of the `INum` interface we attempted earlier:

```
public concept CNum<T>
{
    T Add(T x, T y);
    T Sub(T x, T y);
    T Mul(T x, T y);
    T Abs(T x);
    T Signum(T x);
    T FromInteger(int x);
}
```

Concepts, like interfaces and subtyping relations, are constraint on types. However, we constrain types with concepts in a different way from usual. In Concept-C#, we make generic methods and classes dependent on concepts by adding implicit type parameters, which carry the concept requirement through a C# constraint. We can use `CNum` as follows:

```
public A SomePoly<A, implicit NumA>(A x, A c)
    where NumA : CNum<A> =>
    NumT.Add(NumT.Add(NumT.Mul(x, x), x), c);
```

Formally, when we introduce the `implicit NumA` type parameter, we create a proof obligation that `CNum<T>` holds before we get to use `SomePoly` for `T`. We can write this as a proof tree:

$$\frac{T \vdash \text{CNum} < A > \quad T \vdash A}{T \vdash \text{SomePoly} < A >}$$

To use a concept, we give evidence that for some refinement of the concept type parameters (in this case, `T`), an implementation of the functions in the concept body exists. We do this with an instance declaration. For example, `int` has the following instance!

```
public instance Num_Int : CNum<int>
{
    int Add(int x, int y) => x + y;
    int Sub(int x, int y) => x - y;
    int Mul(int x, int y) => x * y;
    int Abs(int x)         => Math.Abs(x);
    int Signum(int x)      => Math.Sign(x);
}
```

```

    int FromInteger(int x) => x;
}

```

Unlike interface implementations, we can make instances for types that already exist—including primitive types like `int`.

Formally, instances correspond to deduction steps in concept proof trees:

$$\frac{\overline{T \vdash \text{int}}}{T \vdash \text{CNum} < \text{int} >} \text{Num_Int}$$

In most languages with concepts—including Concept-C#—the compiler solves these obligations through type inference. We discuss Concept-C#'s concept inference algorithm later.

Derived instances

Suppose we want to perform `CNum` operations on values that are slow to compute and not guaranteed to be used. It makes sense to make `CNum` instances over `Lazy`—but making a new `Lazy<T>` for each `T` with which we want to do lazy calculations scales poorly.

In Concept-C#, we can give instances implicit type parameters. This means we can make instances conditional on other instances existing. We call such instances *derived*. A derived instance for any `Lazy<A>` for which we have an instance for `A` may look like this:

```

public instance Num_Lazy<A, implicit NA> : CNum<Lazy<A>>
    where NA : CNum<A>
{
    Lazy<A> Add(Lazy<A> x, Lazy<A> y) =>
        new Lazy<A>(() => Add(x.Value, y.Value));
    Lazy<A> Sub(Lazy<A> x, Lazy<A> y) =>
        new Lazy<A>(() => Sub(x.Value, y.Value));
    Lazy<A> Mul(Lazy<A> x, Lazy<A> y) =>
        new Lazy<A>(() => Mul(x.Value, y.Value));
    Lazy<A> Abs(Lazy<A> x) =>
        new Lazy<A>(() => Abs(x.Value));
    Lazy<A> Signum(Lazy<A> x) =>
        new Lazy<A>(() => Signum(x.Value));
    Lazy<A> FromInteger(int x) =>
        new Lazy<A>(() => FromInteger(x));
}

```

Derived instances correspond to intermediate stages in a proof tree:

$$\frac{\overline{\text{CNum} < \text{int} >} \text{Num_Int}}{\text{CNum} < \text{Lazy} < \text{int} > >} \text{Num_Lazy}$$

Operator overloading

Writing numeric code in an functional style, like we did above, is tedious and quickly marches off the right margin of one's text editor. To solve this, we extend C#'s operator overloading support to concepts.

In Concept-C#, concepts can define their own operators. We can rewrite `CNum` to use operators:

```
public concept CNum<T>
{
    T operator +(T x, T y);
    T operator -(T x, T y);
    T operator *(T x, T y);
    T Abs(T x);
    T Signum(T x);
    T FromInteger(int x);
}
```

Concept operators become available for operator overloading if no other operator or operator overload is available, and a suitable concept instance is in scope. When we use `+`, `-`, or `*` on two values of type `T`, there is not already a valid operator or operator overload, and some `CNum<T>` is available, Concept-C# will pick up its definitions.

With concept operator overloading, we can rewrite `SomePoly` to this:

```
public T SomePoly<T, NumT>(T x, T c) where NumT : CNum<T> =>
    x * x + x + c; // NumT is in scope, so use its operators
```

Concept extension methods

Operator overloads help us write concept-based generic arithmetic in a highly idiomatic way. However, if we want to take the absolute value or sign of a `Num<T>`, we still have to write awkward, functional code:

```
var abs = NumT.Abs(x);
var sig = NumT.Signum(x);
```

Ideally, we would like to call `Abs` and `Signum` on `x`, as if they were actual methods of the `T` class.

Just like we extended operator overloads to concepts, we extend extension methods: by prefixing the first parameter of a concept method with `this`, the method becomes a *concept extension method*. Applying this change to `Abs` and `Signum` means we can refactor:

```
var abs = x.Abs();
var sig = x.Signum();
```

Autofilling

At this stage, our `Num_Int` instance looks as follows:

```
public instance Num_Int : CNum<int>
{
    int operator +(int x, int y) => x + y;
    int operator -(int x, int y) => x - y;
    int operator *(int x, int y) => x * y;
    int Abs(this int x)          => Math.Abs(x);
    int Signum(this int x)       => Math.Sign(x);
    int FromInteger(int x)       => x;
}
```

The first three instances are tedious to write, and likely to introduce the dreaded copy-paste error when implemented:

```
int operator -(int x, int y) => x + y; // oops!
```

Thankfully, Concept-C# can infer trivial parts of instances. If an instance method is an operator overload, and there already exists a valid statically defined operator overload or builtin operator for the same types, Concept-C# will fill in the obvious definition of that operator. The above is, thus, equivalent to⁵:

```
public instance Num_Int : CNum<int>
{
    int Abs(this int x)          => Math.Abs(x);
    int Signum(this int x)       => Math.Sign(x);
    int FromInteger(int x)       => x;
}
```

Similarly, suppose we implement `CNum` for a class with methods already called `Abs` and `Signum`. Concept-C# can automatically forward concept extension methods to instance methods on the same class, saving us from needing to write⁶

```
int Abs(thisClazz x) => x.Abs();
```

Defaults

Another operation we often do with numeric types is negation. We left this out of `CNum` since it can be defined entirely in terms of `-` and `FromInteger`. We may want to add it for two reasons: first, most numeric types will let us implement negation in a more performant way, and second, for some types it may be easier to define `-` in terms of `+`

⁵Future versions of Concept-C# will probably prevent the programmer from overloading builtin operators anyway, to prevent surprises.

⁶We don't currently forward static methods, but maybe we should.

and negation. However, this means that we must implement negation in all of our `CNum` instances, even if it is just `x => FromInteger(0) - x`.

Concept-C# allows concepts to provide bodies for methods they specify. These bodies are then used as the *default implementation* of the method in any instance that doesn't either directly implement the method, or implicitly forward it based on the rules above.

Multi-parameter concepts

C# supports generic enumeration of collections through the `IEnumerable<T>` interface, where `T` is the type of elements in the collection. The interface looks like this:

```
public interface IEnumerable<out T> : IEnumerable
{
    IEnumerator<T> GetEnumerator();
}
public interface IEnumerator<out T> : IDisposable, IEnumerator
{
    T Current;
    void Dispose();
    bool MoveNext();
    void Reset();
}
```

If we have a collection `xs` whose type implements this interface, we can use C#'s `foreach` syntax, which looks like this:

```
foreach (var x in xs)
{
    DoSomething(x);
}
```

This pattern also appears at the core of *LINQ to Objects*, where operators such as `select` and `where` transform enumerables into other enumerables. An example of LINQ is:

```
var ys = from x in xs where x > 5 select x * 2;
```

This setup works well—if the author of `xs`'s type had the foresight and patience to implement `IEnumerable<T>`, `IEnumerator<T>`, and the non-generic equivalents `IEnumerable` and `IEnumerator`. We lose any compile-time type information about the `IEnumerator` of `xs` is, which hits performance both by introducing virtual calls and preventing compile-time specialisation. Worse, `IEnumerable<T>` is a hassle to implement properly. Indeed, C# is flexible in the case of `foreach`, and accepts types that implement the correct methods, but not the full interface.

For a satisfying concept implementation of enumerables, we must make the concept generic not only on the type of the collection, but also the type of the enumerator.

Similarly, the concept for enumerators must be generic both on the enumerator and the element. This needs a more general idea of concepts than previously seen. C# already has multi-parameter interfaces and structs, so our implementation of multi-parameter concepts is simple.

We can implement enumerators and enumerables in Concept-C# as follows:

```
public concept CEnumerator<TEnum, TElem>
{
    TElem Current(ref TEnum e);
    void Dispose(ref TEnum e);
    bool MoveNext(ref TEnum e);
    void Reset(ref TEnum e);
}
public concept CEnumerable<TColl, TEnum, TElem>
: CEnumerator<TEnum, TElem>
{
    TEnum GetEnumerator(TColl c);
}
```

When we ask for the enumerator of a CEnumerable, we know *at compile-time* not only that it satisfies CEnumerator, but also its concrete type. We sacrifice some of the abstraction of the interface version, but get possibilities to optimise performance.

Lightening the load

Implementing CEnumerable directly is still overkill for many collections. For example, collections that contain a bounded, ordered number of elements with $O(1)$ lookup given a position in the collection have trivial enumerators, which must be wired up separately for each instance.

Thankfully, derived instances mean we can break down the implementation of a large concept by using smaller concepts. So, if we have the following concept for looking up items in a collection by their index:

```
public concept CIndexable<TColl, TElem>
{
    TElem At(TColl c, int i);
}
```

And the following concept for looking up the length of a collection:

```
public concept CCountable<TColl>
{
    int Count(TColl c);
}
```

We can create a generalised CEnumerable instance capturing the standard pattern of implementing enumerables for these types of collection, which looks as follows:

```
public struct ICEnumerator<TColl>
{
    public TColl src;
    public int pos;
    public int len;
}

public instance Enumerable_IC<TColl, TElem, implicit I, implicit C>
    : CEnumerable<TColl, ICEnumerator<TColl, TElem>, TElem>
    where I : CIndexable<TColl> where C : CCountable<TColl>
{
    ICEnumerator<TColl, TElem> GetEnumerator(TColl c) =>
        new ICEnumerator<TColl, TElem>
            { src = c, pos = 0, len = C.Count(c) };

    TElem Current(ref ICEnumerator<TColl, TElem> e) => e.src[e.pos];

    void Dispose(ref ICEnumerator<TColl, TElem> e) => {};

    bool MoveNext(ref ICEnumerator<TColl, TElem> e)
    {
        if (e.len <= e.pos) return false;
        e.pos++;
        return (e.pos < e.len);
    };

    void Reset(ref ICEnumerator<TColl, TElem> e) => e.pos = -1;
}
```

Now, type creators can write instances for the much simpler CCountable and CIndexable concepts, and the CEnumerable instance (and, through inheritance, the CEnumerator instance) write themselves. For example, we can define list enumeration as possible:

```
public instance Indexable_List<TElem> : CIndexable<List<TElem>>
{
    TElem At(List<TElem> c, int i) => c[i];
}

public instance Countable_List<TElem> : CCountable<List<TElem>>
{
    int Count(List<TElem> c) => c.Count;
}
```

Specialisation

We now have an instance for `CEnumerable<List<T>, ..., T>` without actually implementing `CEnumerable` itself. What if, for performance reasons, we instead want to create a bespoke `CEnumerable` instance for `List<T>`? In fact, we might want to use the same enumerator `List<T>` already exposes—`List<T>.Enumerator`—instead of making our own.

We can write a working instance over `List<T>.Enumerator` as follows:

```
public instance Enumerable_List<TElem>
    : CEnumerable<List<TElem>, List<TElem>.Enumerator, TElem>
{
    List<TElem>.Enumerator GetEnumerator(List<TElem> c) =>
        c.GetEnumerator();

    TElem Current(ref List<TElem>.Enumerator e) => e.Current;
    void Dispose(ref List<TElem>.Enumerator e) => e.Dispose();
    bool MoveNext(ref List<TElem>.Enumerator e) => e.MoveNext();
    void Reset(ref List<TElem>.Enumerator e) => e.Reset();
}
```

This poses a problem: there are now two suitable instances for the concept `CEnumerable<List<TElem>>`. By default, Concept-C# refuses to infer an instance when this happens. This is to help programmers diagnose unintentional overlaps in their instances.

We can still use lists wherever we need `CEnumerable`, by manually specifying which instance we want to use. For example,

```
var general =
    CEnumerable<List<int>,
        Enumerable_IC<List<int>,
            int,
            Indexable_List<int>,
            Countable_List<int>>>
        .GetEnumerator(xs);
var specific =
    CEnumerable<List<int>,
        Enumerable_List<int>>.GetEnumerator(xs);
```

In most cases, we want Concept-C# to pick our more specialised instance over the general one, and don't want to write out the entire instance. Concept-C# has heuristics for *tie-breaking* when multiple instances are available. To use them, we must tell it that the less suitable instance can be *overlapped* by any more suitable instance, or the more

suitable instance can be *overlapping* less suitable instances⁷.

We can turn tie-breaking on by making the following change to our general instance:

```
[Overlappable] // Enumerable_List can overlap this
public instance Enumerable_IC<TColl, TElem, implicit I, implicit C>
    : CEnumerable<TColl, ICEnumerator<TColl, TElem>, TElem>
    where I : CIndexable<TColl> where C : CCountable<TColl>
```

Now, concept inference will pick `Enumerable_List`, if it is in scope.

From interfaces to concepts, and back again

Existing C# code only uses interfaces, not concepts. As a result, it would be hard to use concepts for common patterns such as enumerables if both systems were incompatible.

Adapting interfaces to concepts is as easy as creating a catch-all instance for types implementing that interface, using C#'s existing support for interface constraints. We can port `IEnumerable` code to `CEnumerable` as follows:

```
[Overlappable]
public instance Enumerable_IEnumerable<TColl, TElem>
    : CEnumerable<TColl, TElem, IEnumerator<TElem>>
    where TColl : IEnumerable<TElem>
{
    IEnumerator<TElem> GetEnumerator(TColl c) => c.GetEnumerator();
    void Reset(ref IEnumerator<TElem> e) => e.Reset();
    bool MoveNext(ref IEnumerator<TElem> e) => e.MoveNext();
    TElem Current(ref IEnumerator<TElem> e) => e.Current;
    void Dispose(ref IEnumerator<TElem> e) => e.Dispose();
}
```

This pattern generalises to other concepts. We can port `IComparable` code to `CComparable` as follows:

```
[Overlappable]
public instance Comparable_IComparable<TLhs, TRhs>
    : CComparable<TLhs, TRhs>
    where TLhs : IComparable<TRhs>
{
    int Compare(TLhs l, TRhs r) => l.CompareTo(r);
}
```

Going the other way—creating an `IEnumerable` from `CEnumerable`—is harder, as concepts have no tangible type representation like interfaces do. Interfaces implicitly ‘box’

⁷This system is based on the Glasgow Haskell Compiler's *overlapping instances* extension.

their actual implementation, so we must do a similar job with the concept implementation. This works as follows:

```
class EnumeratorShim<TElem, TState, implicit N> : IEnumerator<TElem>
    where N : CEnumerator<TElem, TState>
{
    private TState _state;

    public EnumeratorShim(TState state)
    {
        _state = state;
    }

    public E Current => N.Current(ref _state);
    object IEnumerator.Current => N.Current(ref _state);
    public bool MoveNext() => N.MoveNext(ref _state);
    public void Reset() { N.Reset(ref _state); }
    void IDisposable.Dispose() { N.Dispose(ref _state); }
}
```

Associated types

Multi-parameter concepts can be hard to infer, as some of the type parameters may not correspond to parameters of the functions we want to call through them. However, it is often the case that one or more of those parameters can be determined precisely from the instance we choose to fulfil the concept.

In the `CEnumerable` example, `TElem` is one such type: it appears only as a return type, and C# cannot infer return types. `TEnum` can also be tricky to work out if we enter the concept through a call to `GetEnumerator`, for the same reason.

Languages that implement multi-parameter concepts have various solutions to this, such as Haskell's type families and functional dependencies, and Rust's associated types. In Concept-C#, we currently have a simple system approximating associated types. Type parameters marked `[AssociatedType]` will be filled in by the concept witness inferer as soon as it fixes an instance for a concept constraint that mentions that type.

We can redefine our `CEnumerator` and `CEnumerable` concepts as follows:

```
public concept CEnumerator<TElem, [AssociatedType] TEnum>
{ /* ... */ }

public concept CEnumerable<TColl,
    [AssociatedType] TElem,
    [AssociatedType] TEnum>
```



```
: CEnumerator<TElem, TEnum>
{ /* ... */ }
```

Now, whenever we try to infer an instance for `CEnumerator` or `CEnumerable`, we will pick up the `TElem` and `TEnum` of the inferred instance.

Standalone instances

Concepts have access to a large amount of features normal extension methods don't have. However, it is possible to use instances as a more powerful alternative to standard extension classes.

By creating a *standalone instance*—an instance that inherits from no concepts—one can attach concept extension methods and operator overloads⁸ to a type without creating a concept to hold them.

For example,

```
public instance GooPlus
{
    public Goo operator +(Goo x, Goo y) => /* ... */
}
```

will add an overload for `+` to all `Goo` as long as `GooPlus` is in scope.

Example: a self-specialising *LINQ to Objects*

TODO

Implementation

Our proposal needs no changes to the Common Language Runtime. We need only modest changes to the compiler and an extension to the core library.

Inference

Our key change is to extend C#'s type inference algorithm in two ways:

- Infer missing concept-witness and associated-type parameters with a new, Haskell-style inference strategy;

⁸It might be a good idea to constrain operator overloads in standalone instances, eg. they can't be used unless there's a concept constraint, or by adding type parameter restrictions.

- Allow partially specified type arguments when the missing arguments are concept witnesses or associated types, and *part-infer* the rest;

**** TODO ****

Related Work

TODO

***Shapes* proposal**

TODO

Conclusions

TODO

Future work

TODO